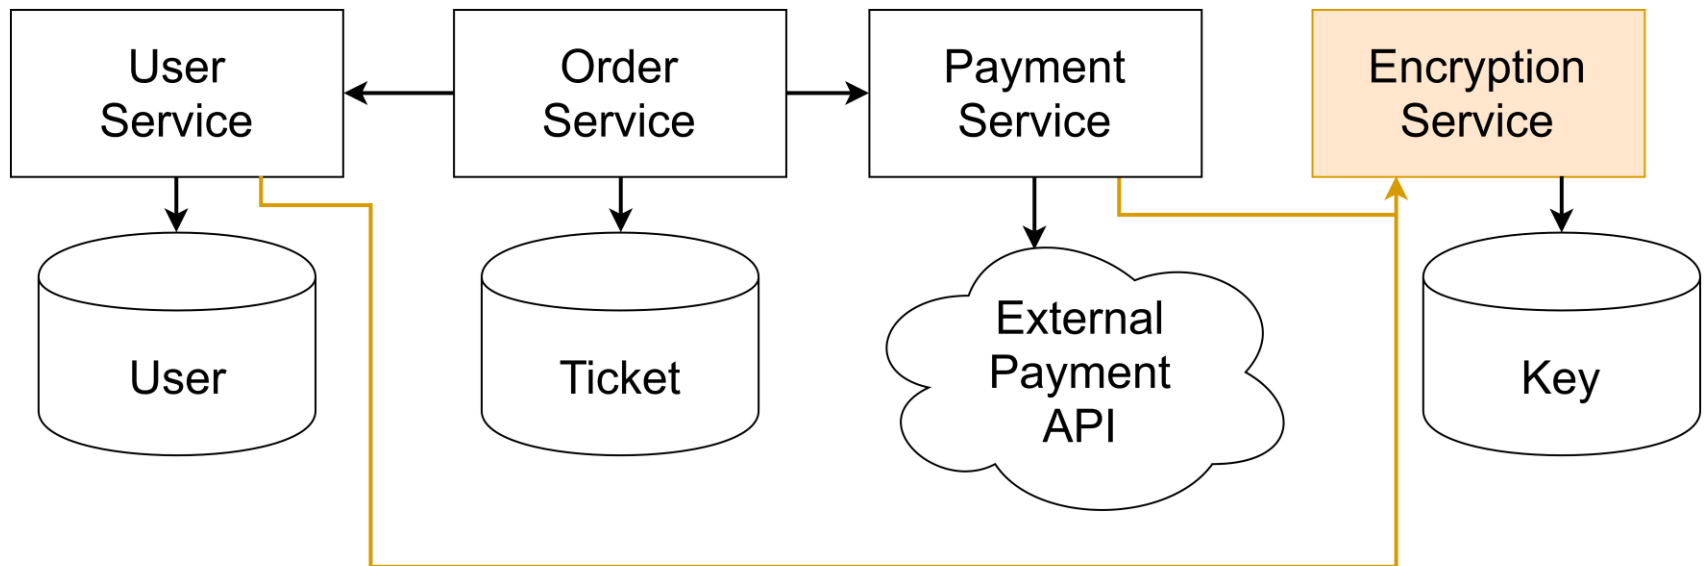


Cloud-based Online Concert Ticketing System

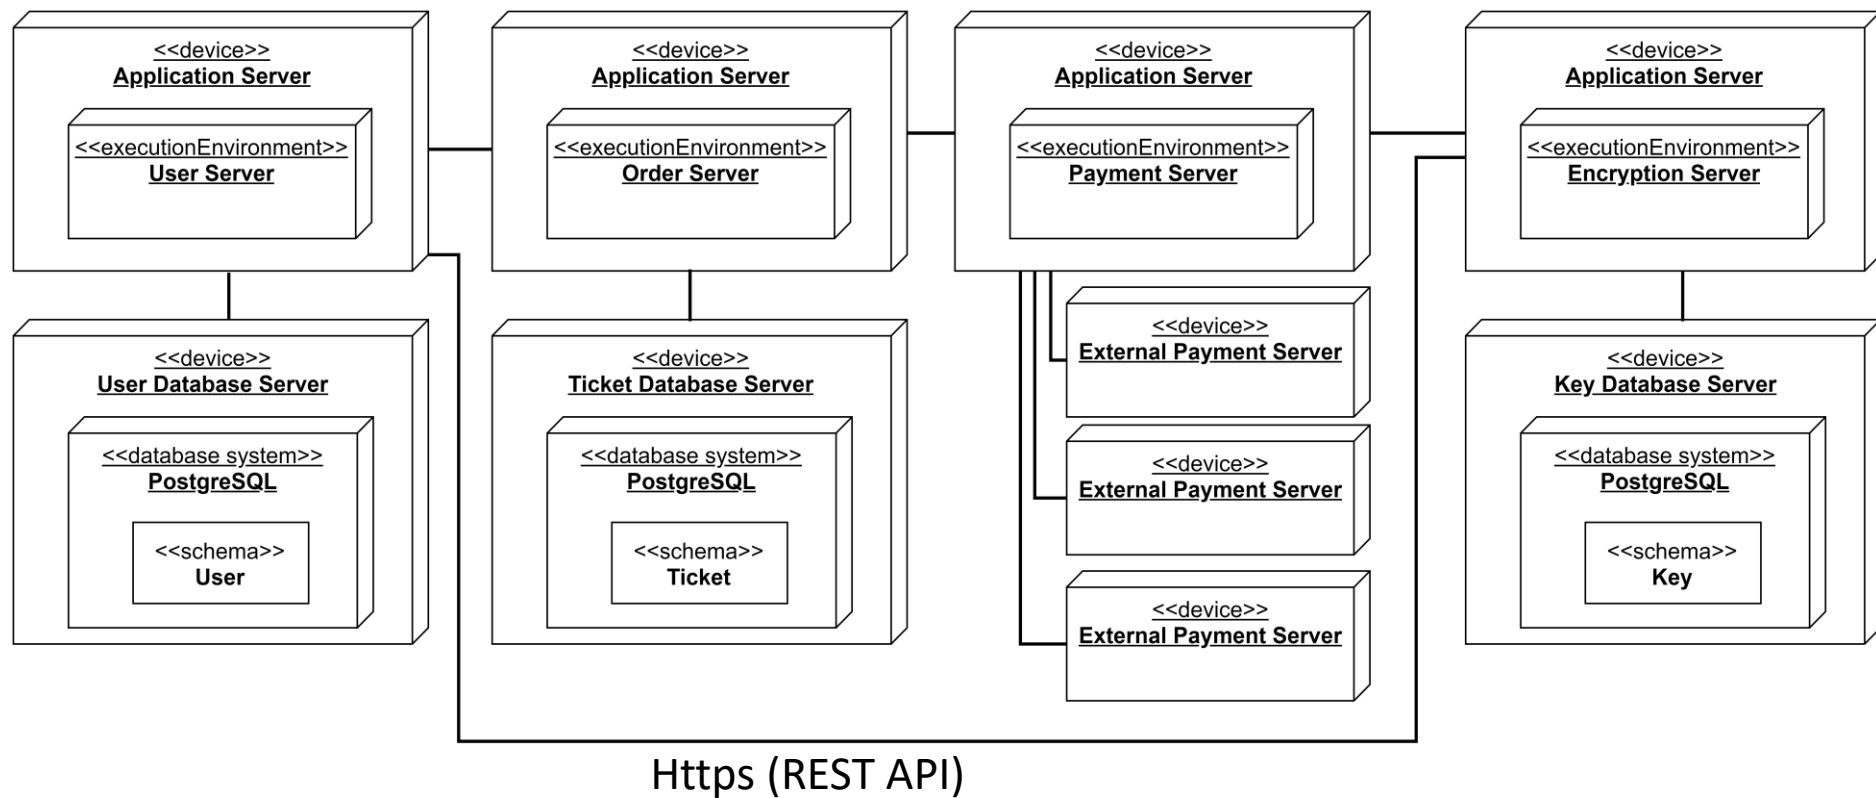
Zhong Xi Lu

Architecture Overview

General Overview



Deployment Diagram



Class Diagram and Database Schemas

UserService
POST users PUT users GET users/{user_id} GET users POST users/verify

User
id username password gender token country city zip_code street card_type card_holder_name card_number expiration_date_month expiration_date_year cvv

OrderService
POST orders

Ticket- <i>i</i>
id user_id token

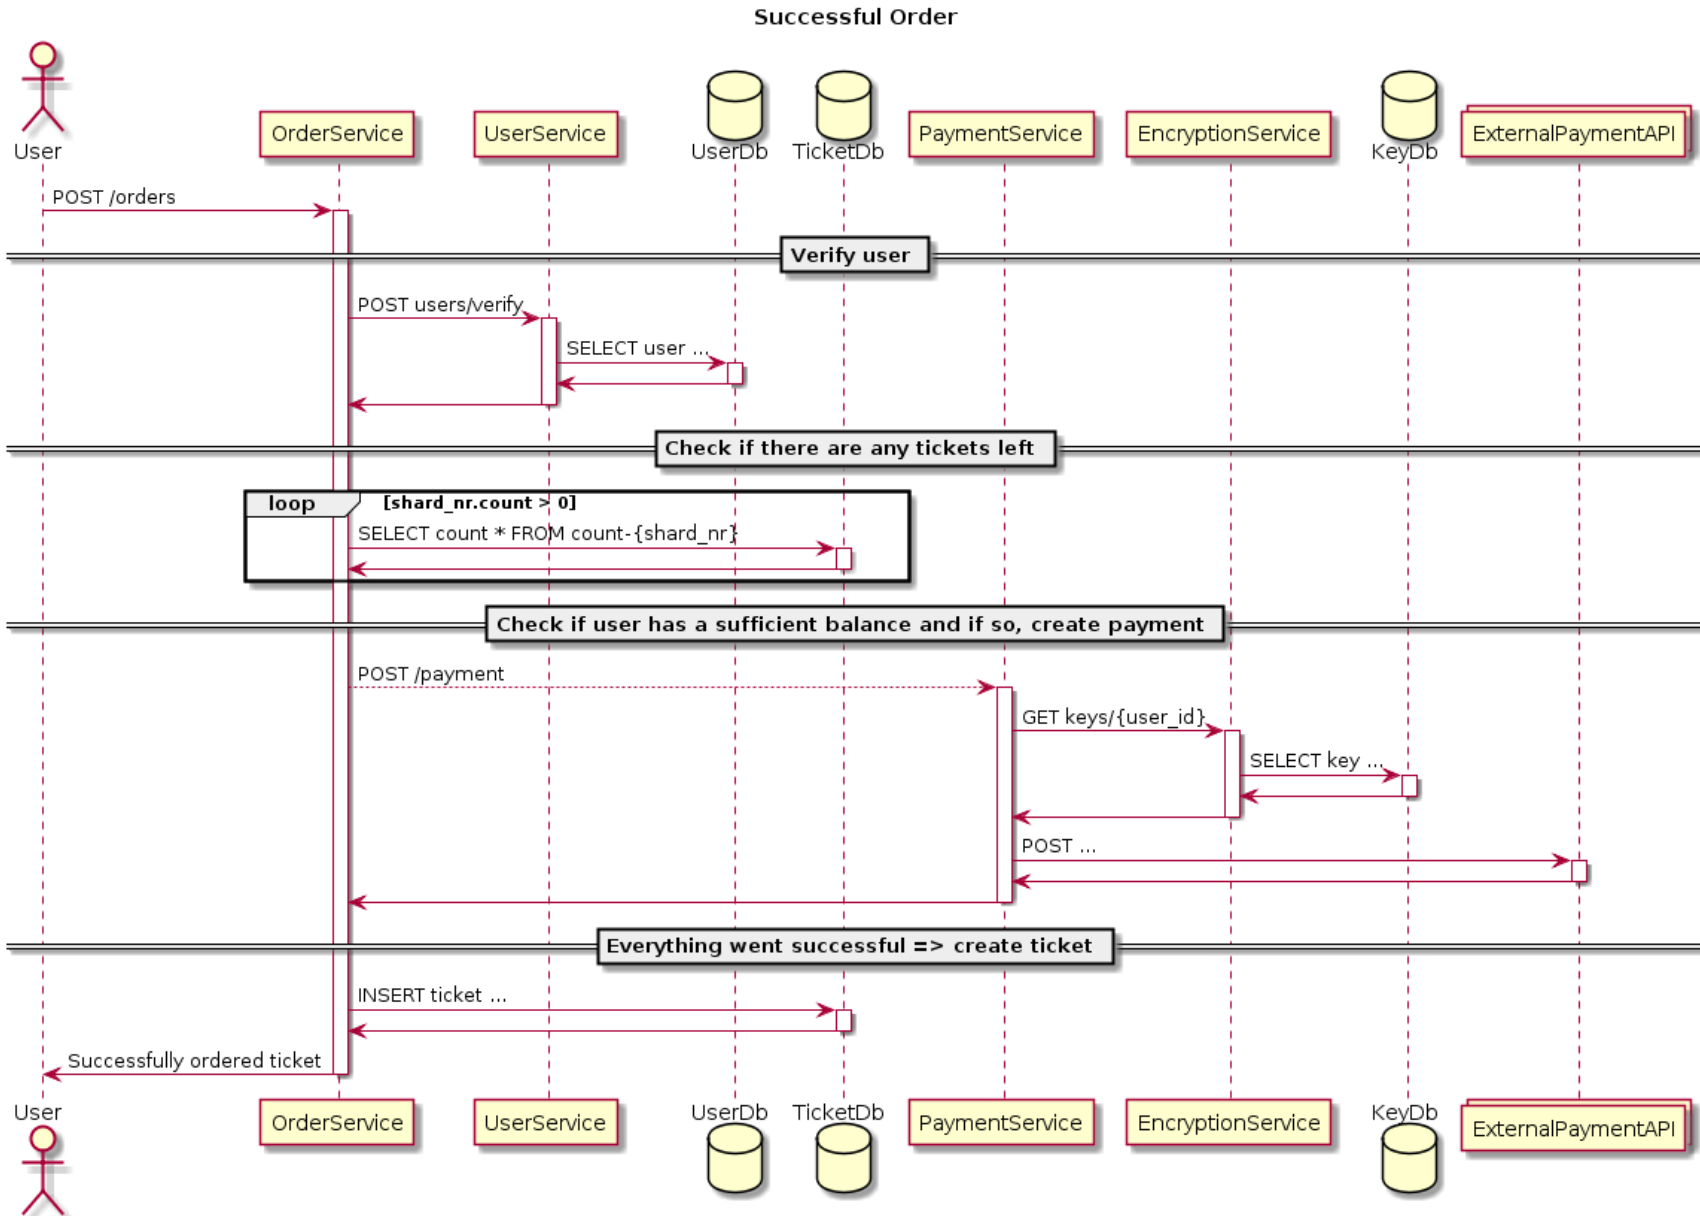
TicketsLeft- <i>i</i>
count

PaymentService
POST payment

EncryptionService
POST keys GET keys/{user_id}

Key
id user_id key IV


Sequence Diagram



Revisiting Old Conclusions


Bottlenecks

- Database => database sharding



- External payment API => solution?


Prediction number of replicas



Prediction number of replicas

- 35.000 replicas?
- Probably only **350** (or more?) replicas = 350 shards
- Maybe even less, one replica can handle **multiple requests concurrently**
- Note: replica \neq server

Prediction response times

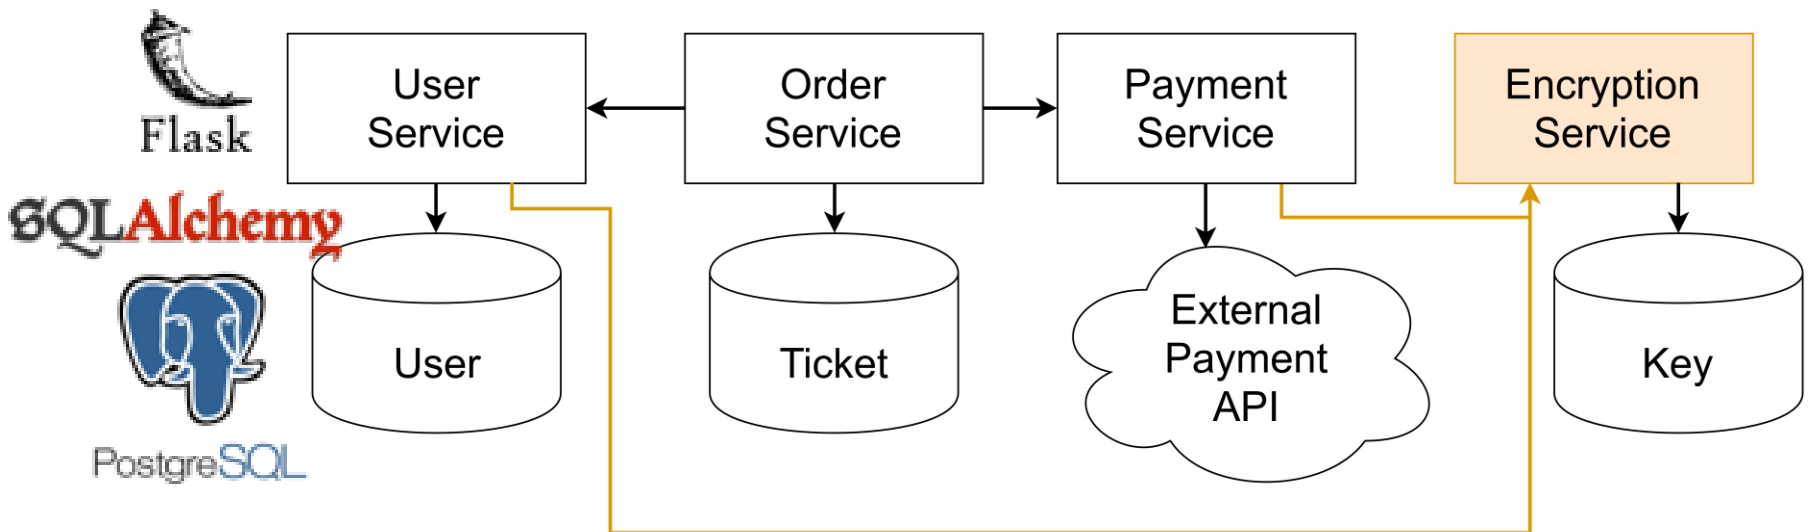


Ordering 350.000 tickets at same time => longest response time = 1.874.549ms \approx 31min.

Ordering **106.300 tickets** at same time => longest response time \approx 600.000ms = **10 min.**

Set up


Architecture




Docker containers




*Note: Everything ran on just **one machine***




User Service




Order Service




Payment Service




Encryption Service



User Database



Ticket Database



Key Database


Autoscaler

- Kubernetes **Horizontal Pod Autoscaler (HPA)**

```
desiredReplicas =  
    ceil[currentReplicas * ( currentMetricValue / desiredMetricValue )]
```

(<https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/#algorithm-details>)

- However, best **scale to maximum at start**
- Scale Order Service



Order Service

Minikube Setup

```
zhongxilu:CapitaSelectaSE$ kubectl get pods --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	encryption-dc8bd8f5-brh9v	1/1	Running	0	6m52s
default	key-db-6dfdcfb5c5d-jt259	1/1	Running	0	6m52s
default	order-5dcb486775-2xjzv	1/1	Running	0	24s
default	order-5dcb486775-g5vlf	1/1	Running	0	24s
default	order-5dcb486775-gv9qz	1/1	Running	0	24s
default	order-5dcb486775-jdzqf	1/1	Running	0	24s
default	order-5dcb486775-qc92r	1/1	Running	0	24s
default	order-5dcb486775-qfp4r	1/1	Running	0	6m52s
default	order-5dcb486775-rp7x8	1/1	Running	0	24s
default	payment-85684648b-5t9sh	1/1	Running	0	6m52s
default	ticket-db-776c9f4459-cjmzr	1/1	Running	0	6m52s
default	user-6d84fc6cdf-dj2d2	1/1	Running	0	6m52s
default	user-db-76b9cdcc9d-lqsw9	1/1	Running	0	6m52s
kube-system	coredns-584795fc57-p52m7	1/1	Running	4	9m19s
kube-system	coredns-584795fc57-trhdx	1/1	Running	4	9m19s
kube-system	etcd-minikube	1/1	Running	0	8m17s
kube-system	kube-addon-manager-minikube	1/1	Running	0	8m7s
kube-system	kube-apiserver-minikube	1/1	Running	0	8m10s
kube-system	kube-controller-manager-minikube	1/1	Running	0	8m19s
kube-system	kube-proxy-gplvg	1/1	Running	0	9m19s
kube-system	kube-scheduler-minikube	1/1	Running	0	8m23s
kube-system	storage-provisioner	1/1	Running	0	9m16s


Results

Load Test Script

- Before: n registered users
- Retrieve all users
- Order ticket for each registered user at the same time (**n requests**)


Number of replicas for Order service

Response times for different #replicas




Ordering 10.000 tickets

Number of shards



Ordering 1000 tickets

External API bottleneck



Ordering 1000 tickets

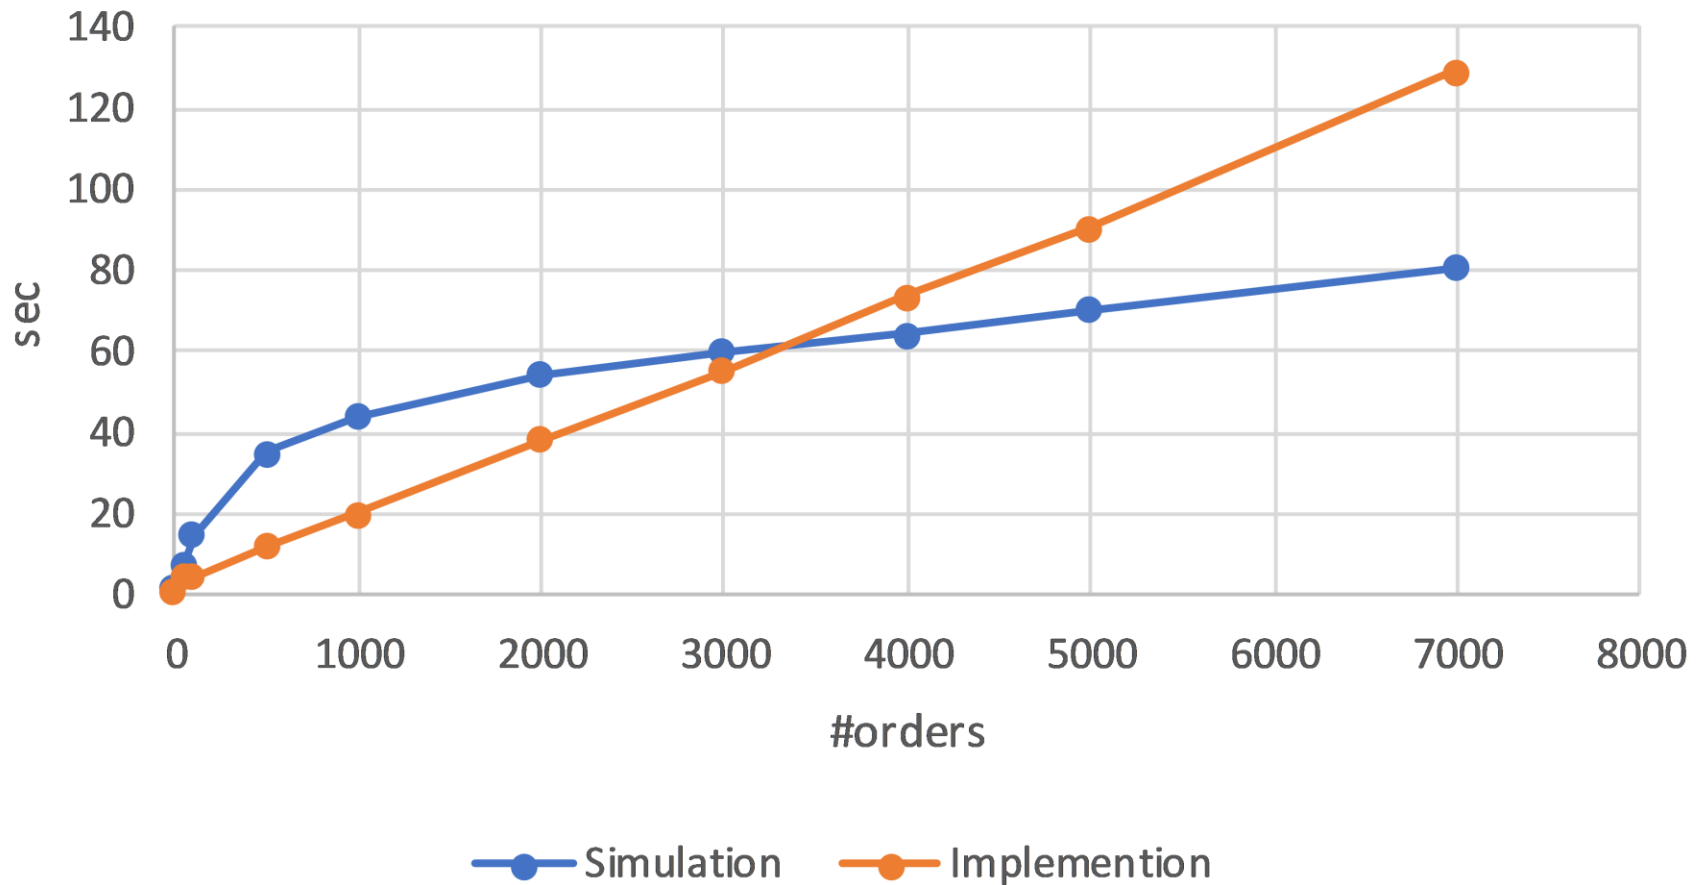
Encryption

- Ordering 10.000 tickets at the same time:

	w/ encryptions	w/o encryptions
Shortest response time	4.5 sec	4.29 sec
Average response time	90.57 sec	87.711 sec
Longest response time	178.023 sec	172.408 sec

Comparison with ABS simulations

Response times for different #orders



Peak Load

- **31.186 tickets** ordered at same time for response time less than 10 min
- *Important:* only using **one machine**

Conclusion (again)

- So does it respond in less than 10 min during peak load?
 - *It depends...*
 - *How much is the actual **peak load**?*
 - *How much **resources** available?*
 - *Processing time **external payment API**?*
 - But system **scales dynamically** on cloud